

Stanford's Graph-based Neural Dependency Parser at the CoNLL 2017 Shared Task

Timothy Dozat Peng Qi Christopher D. Manning

Stanford University

August 6, 2017

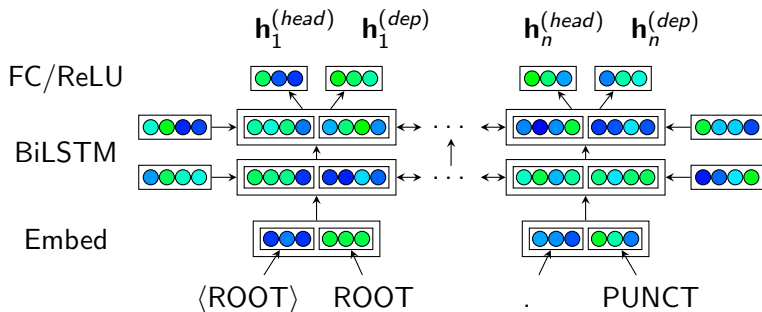


Overview

- 1 Parser
- 2 UPOS/XPOS tagger
- 3 Character-level embedding model
- 4 Results
- 5 Noteworthy hyperparameters

Overview: Architecture


Almost everything builds on this structure (cf. Dozat and Manning (2017)):



Parser

Unlabeled parser: LSTM

- Bidirectional LSTM over word/tag embeddings (more on embeddings later)
- Two separate FC ReLU layers
 - One representing each token as a dependent trying to find (attend to) its head
 - One representing each token as a head trying to find (be attended to by) its dependents



Unlabeled parser: Self-attention

- Biaffine self-attention layer to score possible heads for each dependent

$$\mathbf{s}_i^{(arc)} = H^{(arc-head)} (W \oplus \mathbf{b} \quad \mathbf{h}_i^{(arc-dep)} \oplus 1)^T$$


- Train with cross-entropy
- Apply a spanning tree algorithm at inference time

Note: This is just an affine layer with a linear transformation!

$$\mathbf{s}_i = H^{(arc-head)} (W \mathbf{h}_i^{(arc-dep)} + \mathbf{b})$$

Labeler: LSTM

- Take the topmost BiLSTM vectors used for the unlabeled parser
- Two more separate FC ReLU layers:
 - One representing each token as a dependent trying to determine its label
 - One representing each token as a head trying to determine its dependents' labels



Labeler: Classifier

- Biaffine layer to score possible relations for each best-head/dependent pair

$$\mathbf{s}_i^{(rel)} \quad \mathbf{h}_{y_i}^{(rel-head)} \oplus 1 \quad \mathbf{U} \quad \mathbf{h}_i^{(rel-dep)} \oplus 1$$

The diagram illustrates the biaffine layer operation. It shows a vector $\mathbf{s}_i^{(rel)}$ (represented by a box with two blue circles) multiplied by a vector $\mathbf{h}_{y_i}^{(rel-head)} \oplus 1$ (represented by a box with three green circles). This is followed by a matrix \mathbf{U} (represented by a 4x4 grid of blue and green circles) and another vector $\mathbf{h}_i^{(rel-dep)} \oplus 1$ (represented by a box with three blue circles). The result is a scalar value (represented by a box with one green circle).

- Train with softmax cross-entropy, added to the loss of the unlabeled parser

Note: this is just a linear model with interaction effects!

$$\text{label.scores} \sim \text{head.state} * \text{dep.state}$$


Tagger

Tagger: Motivation

- **Problem:** Dozat and Manning's (2017) parser had lower label accuracy than we wanted
- **Idea:** Better POS tag quality might improve label score
- **Question:** Will improved POS tag accuracy result in better parsers?

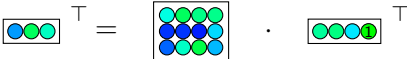
Tagger: LSTM

- BiLSTM (distinct from parser BiLSTM!) over word embeddings
- Two separate FC ReLU layers:
 - One for UPOS tags
 - One for XPOS tags




Tagger: Classifiers

- Affine layers to score possible tags for each word


$$\mathbf{s}_i^{(pos)} = \mathbf{W} \oplus \mathbf{b} \cdot \mathbf{h}_i^{(pos)} \oplus 1$$


- Train jointly by adding together softmax cross-entropy
- When using in the main parser, add UPOS and XPOS embeddings together (eltwise)



Tagger: Experiment

- Systems with our tagger outperformed systems with baseline tagger (Straka et al., 2015) ($p < .05$) or no tagger ($p < .05$)
- Parser performance correlated with tagger performance (ours vs. baseline) ($p < .05$)



Character Model


Character model: Motivation

- **Problem:** Many shared task languages have complex morphology
 - Grammatical functions indicated more by word form than relative location
 - Rare words with highly predictive suffixes won't be attested in the frequent word embedding matrix
 - Extreme sparsity may yield low-quality pretrained embeddings
- **Idea:** Compose word embeddings orthographically with a character-based embedding model
- **Question:** Does this improve accuracy on inflectionally rich languages?

Character model: LSTM


- Unidirectional LSTM over character embeddings
- Concatenate two sources of information:
 - Linear attention over top hidden states (Cao and Rei, 2016)
 - Final cell state (Ballesteros et al., 2015)

Attn/Final Cell




Character model: Embedding

- Linearly transform to the desired size





- When using in the parser/tagger, add with pretrained and frequent-token embeddings (eltwise)



Character model: Experiment

- Systems trained with a character model outperformed models trained without ($p < .05$)
- Improvement correlated with morphological complexity ($p < .05$)



Results

Results

Treebanks	UPOS	XPOS	UAS	LAS	CLAS
All treebanks	93.09	82.27	81.30	76.30	72.57
Large treebanks	95.58	94.56	85.16	81.77	78.40
Parallel treebanks	88.25	30.66	80.17	73.73	69.88
Small treebanks	87.02	82.03	70.19	61.02	54.76
Surprise treebanks	–	–	54.47	40.57	37.41
System	UPOS	XPOS	UAS	LAS	CLAS
Dozat et al.	93.09	82.27	81.30	76.30	72.57
Björkelund et al.	<i>91.98</i>	64.84	79.90	74.42	70.18
Yu et al.	91.00	<i>79.93</i>	74.22	68.41	63.24
Shi et al.	90.88	79.80	<i>80.35</i>	<i>75.00</i>	<i>70.91</i>

Hyperparameters (Protips)

Noteworthy hyperparameters

Dropout

- Lots of dropout: `keep_prob` is .67 throughout the whole network
- Embedding dropout
 - Drop token/tag embeddings independently
 - When one is dropped, the other is scaled up to compensate
 - When both are dropped, replace with zeros
 - Seems to work better than random vector/UNK replacement
- Same-mask recurrent dropout (Gal and Ghahramani, 2016)
 - Drop input connections *and* recurrent connections
 - Drop the same connections at each recurrent timestep
 - Seems to work better than traditional dropout/zoneout (Krueger et al., 2017)

Noteworthy hyperparameters

Adam

- Adam optimizer (Kingma and Ba, 2015) with $\beta_1 = \beta_2 = .9$
- For embedding matrices, only decay \mathbf{m} and \mathbf{v} accumulators for tokens that were used in the minibatch
 - I.e. for words that *are* attested in the minibatch, we apply Adam's accumulator update rule:

$$\mathbf{m}_t = \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t$$

$$\mathbf{v}_t = \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2$$

- But for words that *aren't*, we don't update the accumulators, preventing them from decaying down to zero for uncommon words
- Note: this is not the behavior of most ML toolkits!

Noteworthy hyperparameters

Initialization



- Preference for initializing to zero wherever possible
 - Bias terms
 - Final linear layers (character model, output layers)
 - Word/POS embeddings (other than pretrained)
- Otherwise, we use orthonormal initialization (Saxe et al., 2014)

Recurrent Cells

- LSTMs vastly outperformed GRUs and slightly outperformed coupled input-forget LSTMs (Greff et al., 2016)
- Adding a forget bias hurts performance

Nonprojectivity

- Our system outperforms UDPipe v1.1 (transition-based) by a larger margin on treebanks with many crossing arcs ($p < .05$)
- Stronger correlation for treebanks with more crossing arcs in the test set than in the training set ($p < .05$)



Thanks for listening!

References I

- Ballesteros, M., Dyer, C., and Smith, N. A. (2015). Improved transition-based parsing by modeling characters instead of words with lstms. *EMNLP*.
- Björkelund, A., Falenska, A., Yu, X., and Kuhn, J. (2017). Ims at the conll 2017 ud shared task: Crfs and perceptrons meet neural networks. In *Proceedings of the CoNLL 2017 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies*, pages 40–51, Vancouver, Canada. Association for Computational Linguistics.
- Cao, K. and Rei, M. (2016). A joint model for word embedding and word morphology. *ACL 2016*, page 18.

References II

- Dozat, T. and Manning, C. D. (2017). Deep biaffine attention for neural dependency parsing. *ICLR 2017*.
- Dozat, T., Qi, P., and Manning, C. D. (2017). Stanford's graph-based neural dependency parser at the conll 2017 shared task. In *Proceedings of the CoNLL 2017 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies*, pages 20–30, Vancouver, Canada. Association for Computational Linguistics.
- Gal, Y. and Ghahramani, Z. (2016). Dropout as a bayesian approximation: Representing model uncertainty in deep learning. *International Conference on Machine Learning*.

References III

- Greff, K., Srivastava, R. K., Koutník, J., Steunebrink, B. R., and Schmidhuber, J. (2016). LSTM: A search space odyssey. *IEEE Transactions on Neural Networks and Learning Systems*.
- Kingma, D. and Ba, J. (2015). Adam: A method for stochastic optimization. *International Conference on Learning Representations*.
- Krueger, D., Maharaj, T., Kramár, J., Pezeshki, M., Ballas, N., Ke, N. R., Goyal, A., Bengio, Y., Larochelle, H., Courville, A., et al. (2017). Zoneout: Regularizing rnns by randomly preserving hidden activations. *ICLR 2017*.
- Saxe, A. M., McClelland, J. L., and Ganguli, S. (2014). Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. *ICLR 2014*, abs/1312.6120.

References IV

- Shi, T., Wu, F. G., Chen, X., and Cheng, Y. (2017). Combining global models for parsing universal dependencies. In *Proceedings of the CoNLL 2017 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies*, pages 31–39, Vancouver, Canada. Association for Computational Linguistics.
- Straka, M., Hajic, J., Straková, J., and Hajic jr, J. (2015). Parsing universal dependency treebanks using neural networks and search-based oracle. In *International Workshop on Treebanks and Linguistic Theories (TLT14)*, page 208.

References V

Yu, K., Sofroniev, P., Schill, E., and Hinrichs, E. (2017). The parse is darc and full of errors: Universal dependency parsing with transition-based and graph-based algorithms. In *Proceedings of the CoNLL 2017 Shared Task: Multilingual Parsing from Raw Text to Universal Dependencies*, pages 126–133, Vancouver, Canada. Association for Computational Linguistics.