



Deep Biaffine Attention for Neural Dependency Parsing

Timothy Dozat Christopher D. Manning

Stanford University



Goals


- ▶ Much research has been devoted to developing neural dependency parsers with complex, task-specific architecture
- ▶ Typical approach: use specialized neural networks to predict discrete actions in a dedicated, **transition-based** parsing algorithm **SyntaxNet** AKA **Parsey McParseface** (Andor et al., 2016): Feedforward network with beam search and CRF loss
- ▶ **Ablated RNN Grammar** (Kuncoro et al., 2016): Stack-LSTM with bidirectional LSTM for phrase composition (SOTA)
- ▶ **Can we get competitive (or even superior) parsing results with a simple architecture using general-purpose components?**

Dependency Parsing

- ▶ Automatically annotate sentences, focusing on the functional role each phrase plays
 - ▶ **Head:** Edge source, more contentful role (predicate \rightarrow arguments)
 - ▶ **Dependent:** Edge target
 - ▶ **Label:** Edge type (*Nominal* SUBJECT, *Adjectival* CLause)
-
- ▶ Particularly useful for NLU tasks, such as semantic parsing or knowledge base population
 - ▶ **Graph-based** approach to parsing: assign weights to each possible edge, construct a maximum spanning tree

LSTM

Step one: BiLSTM over the sequence of word and part of speech tag embeddings, take all topmost LSTM states R ($= \text{stack}_{i=1}^n(\mathbf{r}_i)$)



Variable-class classification (= attention)

- ▶ We want to predict heads (classes) given dependents (inputs), but the number of possible heads changes from sentence to sentence
- ▶ Thus, we want to predict $P(y_i^{(edge)} = j | \mathbf{r}_i; \mathbf{r}_j)$
- ▶ $\text{softmax}(RU^{(1)}\mathbf{r}_i + R\mathbf{u}^{(2)})$ achieves this naturally

$$P(j | \mathbf{r}_i; \mathbf{r}_j) \propto \exp(\mathbf{r}_i^\top U^{(1)} \mathbf{r}_j) \exp(\mathbf{u}^{(2)\top} \mathbf{r}_j)$$


$$\bullet = \begin{bmatrix} \text{dep} & \text{head} \end{bmatrix} \cdot \begin{bmatrix} \text{dep} & \text{head} \end{bmatrix}^\top$$

- ▶ After deciding on an edge from j to i , we want to predict the label
 - ▶ This time, we want to predict $P(y_i^{(label)} = l | \mathbf{r}_i, \mathbf{r}_{y_i^{(edge)}})$
 - ▶ We can use $\text{softmax}(\mathbf{r}_{y_i}^\top U^{(1)} \mathbf{r}_i + U^{(2)}(\mathbf{r}_{y_i} \oplus \mathbf{r}_i) + \mathbf{b})$ to model this
- $$P(l | \mathbf{r}_i, \mathbf{r}_{y_i^{(edge)}}) \propto \exp(\mathbf{r}_i^\top U_l^{(1)} \mathbf{r}_{y_i}) \exp(\mathbf{r}_i^\top \mathbf{u}_l^{(2)}) \exp(\mathbf{r}_{y_i}^\top \mathbf{u}_l^{(3)}) \exp(b_l)$$
- ▶ Closely related to linear models with interactions

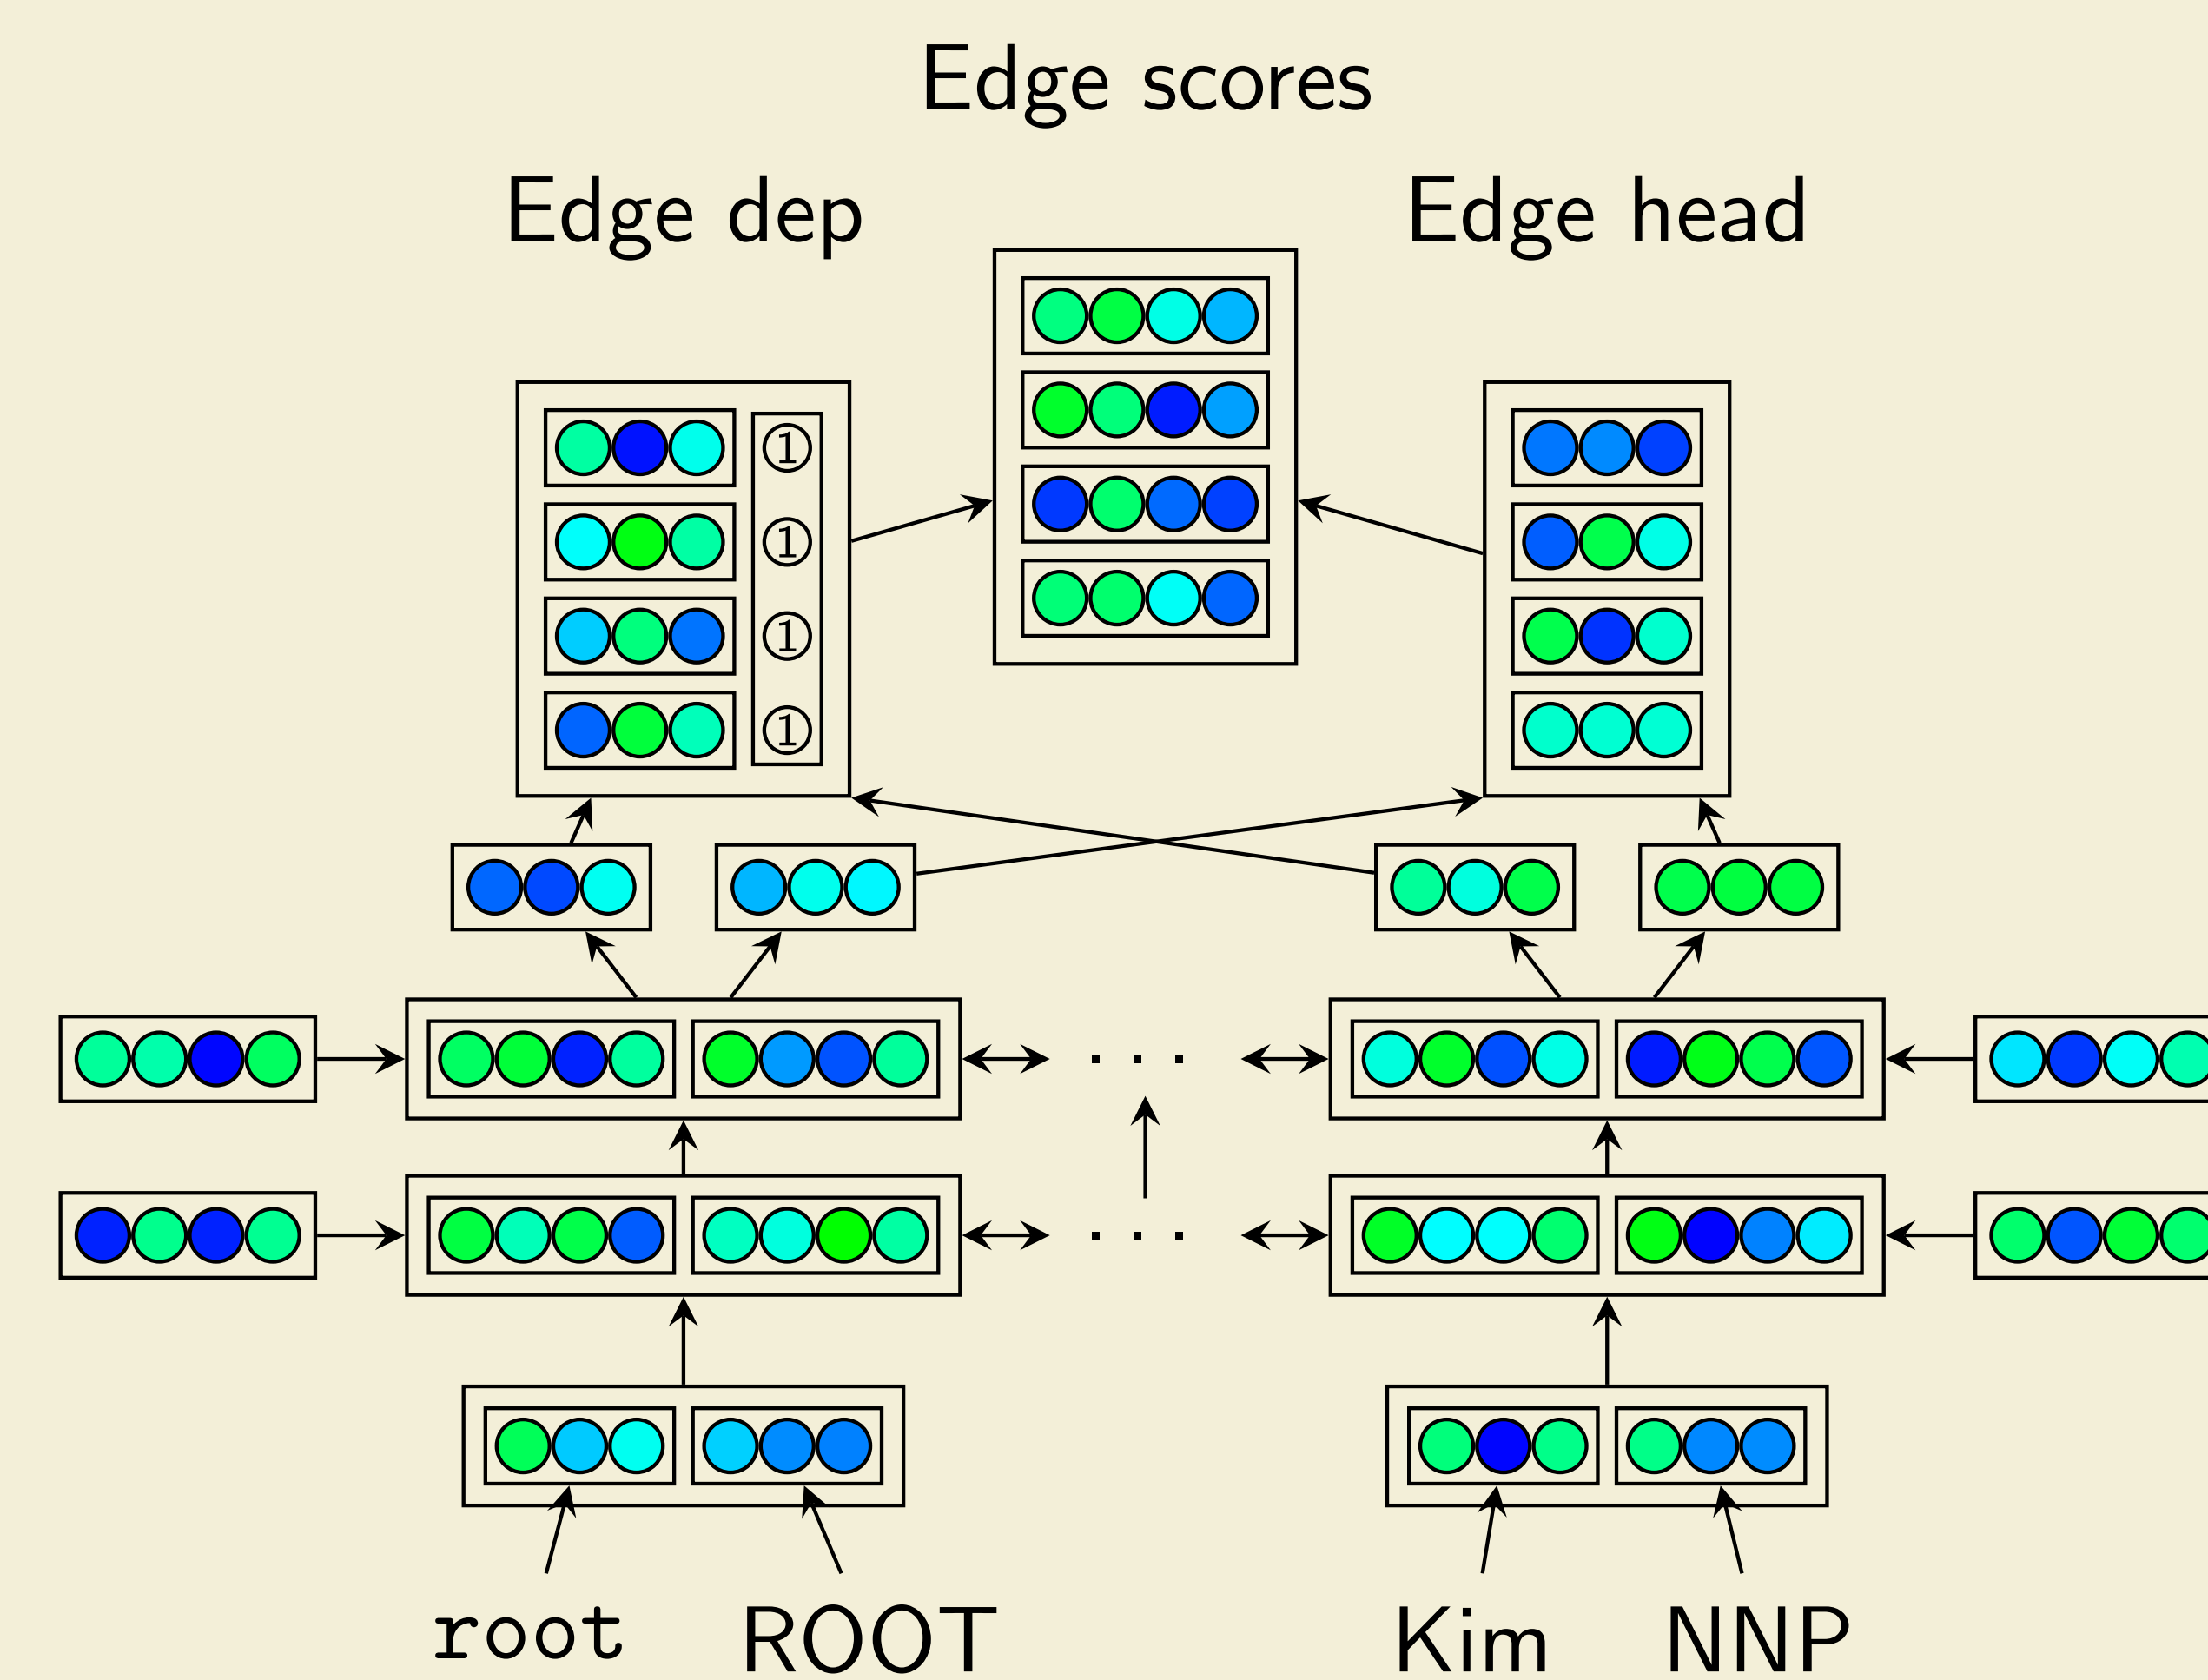
$$\text{scores} \sim \text{head.vector} * \text{dep.vector}$$

Practical modifications

- ▶ Everything is so big!
- ▶ We can get more control over the tradeoffs between speed, overfitting, and underfitting by shrinking \mathbf{r}_i with smaller MLPs before the biaffine output layers (*deep biaffine* model as opposed to *shallow biaffine*)
- ▶ Result: four representations for each word
- ▶ Naturally reflects the intuition that the relationships we want to capture are asymmetric



Final model (edge scorer)



Hyperparameters

| Param | Value | Param | Value |
|----------------|------------------------|--------------------|--------|
| Embedding size | 100 | Embedding dropout | 33% |
| LSTM size | 400 | LSTM dropout | 33% |
| Edge MLP size | 500 | Edge MLP dropout | 33% |
| Label MLP size | 100 | Label MLP dropout | 33% |
| LSTM depth | 3 | MLP depth | 1 |
| α | $2e^{-3}$ | β_1, β_2 | .9 |
| Annealing | $.75^{\frac{t}{5000}}$ | t_{max} | 50,000 |

- ▶ Relatively large network (other models use ~ 100 LSTM dims)
- ▶ Highly regularized with dropout
- ▶ Reducing *Adam's* β_2 from .999 to .9 significantly improved performance ($p < .05$)

Related work

- ▶ Transition-based
 - Nivre et al. (2006):** Feature-based
 - Chen and Manning (2014):** First successful neural parser
 - Andor et al. (2016):** Extend with beam search / CRF loss
 - Kuncoro et al. (2016):** Extend with LSTMs (SOTA)
- ▶ Graph-based
 - McDonald and Pereira (2006):** Feature-based
 - Kiperwasser and Goldberg (2016):** First neural graph-based parser
 - Cheng et al. (2016):** Keep track of previous decisions
 - Hashimoto et al. (2016):** Jointly learn tagging & chunking

PTB Results

| Type | Model | SD 3.3.0 UAS | SD 3.3.0 LAS | CTB UAS | CTB LAS |
|------------|---------------------------------|--------------|--------------|-------------|-------------|
| Transition | Ballesteros et al. (2016) | 93.6 | 91.4 | 87.7 | 86.2 |
| | Andor et al. (2016) | 94.6 | 92.8 | — | — |
| | Kuncoro et al. (2016) | 95.8 | 94.6 | — | — |
| Graph | Kiperwasser and Goldberg (2016) | 93.9 | 91.9 | 87.6 | 86.1 |
| | Cheng et al. (2016) | 94.1 | 91.5 | 88.1 | 85.7 |
| | Hashimoto et al. (2016) | 94.7 | 92.9 | — | — |
| | Deep biaffine | 95.7 | 94.1 | 89.3 | 88.2 |

CoNLL 09 Results

| Model | Catalan UAS | Catalan LAS | Chinese UAS | Chinese LAS | Czech UAS | Czech LAS |
|---------------|-------------|-------------|-------------|-------------|-------------|-------------|
| Andor et al. | 92.7 | 89.8 | 84.7 | 80.9 | 88.9 | 84.6 |
| Deep biaffine | 94.7 | 92.0 | 88.9 | 85.4 | 92.1 | 87.4 |

| Model | English UAS | English LAS | German UAS | German LAS | Spanish UAS | Spanish LAS |
|---------------|-------------|-------------|-------------|-------------|-------------|-------------|
| Andor et al. | 93.2 | 91.2 | 90.9 | 89.2 | 92.6 | 90.0 |
| Deep biaffine | 95.2 | 93.2 | 93.5 | 91.4 | 94.3 | 91.7 |

Affect of classifier type (SD 3.5.0)

| Model | Classifier UAS | Classifier LAS | Sents/sec |
|------------------------------|----------------|----------------|--------------|
| Deep biaffine | 95.8 | 94.2 | 410.9 |
| Shallow biaffine | 95.7 | 94.0* | 299.0 |
| Shallow b. (50% MLP dropout) | 95.7 | 94.1* | 300.1 |
| Shallow b. (300d LSTM) | 95.6* | 93.9* | 373.2 |
| Traditional attention | 95.5* | 93.9* | 367.4 |

(Statistical significances are marked with an asterisk)

Conclusion

- ▶ Our simple, straightforward parser uses only neural components, effectively no task-specific architecture
- ▶ **Substantially outperforms most more complex neural transition-based parsers**
- ▶ **Substantially outperforms all other neural graph-based parsers**
- ▶ The biaffine approach to attention is theoretically justified, here beats the more traditional approach
- ▶ Adding final MLP layers to the LSTM helps to maximize speed and performance, captures head-dependent asymmetries
- ▶ This work provides a fast, simple, high-performing baseline against which to test more complex architectures

References

- Andor, D., Alberti, C., Weiss, D., Severyn, A., Presta, A., Ganchev, K., Petrov, S., and Collins, M. (2016). Globally normalized transition-based neural networks. In *Association for Computational Linguistics*.
- Ballesteros, M., Goldberg, Y., Dyer, C., and Smith, N. A. (2016). Training with exploration improves a greedy stack-LSTM parser. *Proceedings of the conference on empirical methods in natural language processing*.
- Chen, D. and Manning, C. D. (2014). A fast and accurate dependency parser using neural networks. In *Proceedings of the conference on empirical methods in natural language processing*, pages 740–750.
- Cheng, H., Fang, H., He, X., Gao, J., and Deng, L. (2016). Bi-directional attention with agreement for dependency parsing. *arXiv preprint arXiv:1608.02076*.
- Hashimoto, K., Xiong, C., Tsuruoka, Y., and Socher, R. (2016). A joint many-task model: Growing a neural network for multiple nlp tasks. *arXiv preprint arXiv:1611.01587*.
- Kiperwasser, E. and Goldberg, Y. (2016). Simple and accurate dependency parsing using bidirectional LSTM feature representations. *Transactions of the Association for Computational Linguistics*, 4:313–327.
- Kuncoro, A., Ballesteros, M., Kong, L., Dyer, C., Neubig, G., and Smith, N. A. (2016). What do recurrent neural network grammars learn about syntax? *CoRR*, abs/1611.05774.
- McDonald, R. T. and Pereira, F. C. (2006). Online learning of approximate dependency parsing algorithms. In *EACL*.
- Nivre, J., Hall, J., and Nilsson, J. (2006). Maltparser: A data-driven parser-generator for dependency parsing. In *Proceedings of LREC*, volume 6, pages 2216–2219.